

Convolutional Neural Networks

Project: Write an Algorithm for Landmark Classification

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to HTML, all the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Why We're Here

Photo sharing and photo storage services like to have location data for each photo that is uploaded. With the location data, these services can build advanced features, such as automatic suggestion of relevant tags or automatic photo organization, which help provide a compelling user experience. Although a photo's location can often be obtained by looking at the photo's metadata, many photos uploaded to these services will not have location metadata available. This can happen when, for example, the camera capturing the picture does not have GPS or if a photo's metadata is scrubbed due to privacy concerns.

If no location metadata for an image is available, one way to infer the location is to detect and classify a discernible landmark in the image. Given the large number of landmarks across the world and the immense volume of images that are uploaded to photo sharing services, using human judgement to classify these landmarks would not be feasible.

In this notebook, you will take the first steps towards addressing this problem by building models to automatically predict the location of the image based on any landmarks depicted in the image. At the end of this project, your code will accept any user-supplied image as input and suggest the top k most relevant landmarks from 50 possible landmarks from across the world. The image below displays a potential sample output of your finished project.



The Road Ahead

We break the notebook into separate steps. Feel free to use the links below to navigate the notebook.

- [Step 0](#): Download Datasets and Install Python Modules
 - [Step 1](#): Create a CNN to Classify Landmarks (from Scratch)
 - [Step 2](#): Create a CNN to Classify Landmarks (using Transfer Learning)
 - [Step 3](#): Write Your Landmark Prediction Algorithm
-

Step 0: Download Datasets and Install Python Modules

Note: if you are using the Udacity workspace, **YOU CAN SKIP THIS STEP**. The dataset can be found in the `/data` folder and all required Python modules have been installed in the workspace.

Download the [landmark dataset](#). Unzip the folder and place it in this project's home directory, at the location `/landmark_images`.

Install the following Python modules:

- cv2
 - matplotlib
 - numpy
 - PIL
 - torch
 - torchvision
-

Step 1: Create a CNN to Classify Landmarks (from Scratch)

In this step, you will create a CNN that classifies landmarks. You must create your CNN *from scratch* (so, you can't use transfer learning yet!), and you must attain a test accuracy of at least 20%.

Although 20% may seem low at first glance, it seems more reasonable after realizing how difficult of a problem this is. Many times, an image that is taken at a landmark captures a fairly mundane image of an animal or plant, like in the following picture.



Just by looking at that image alone, would you have been able to guess that it was taken at the Haleakalā National Park in Hawaii?

An accuracy of 20% is significantly better than random guessing, which would provide an accuracy of just 2%. In Step 2 of this notebook, you will have the opportunity to greatly improve accuracy by using transfer learning to create a CNN.

Remember that practice is far ahead of theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

(IMPLEMENTATION) Specify Data Loaders for the Landmark Dataset

Use the code cell below to create three separate [data loaders](#): one for training data, one for validation data, and one for test data. Randomly split the images located at `landmark_images/train` to create the train and validation data loaders, and use the images located at `landmark_images/test` to create the test data loader.

All three of your data loaders should be accessible via a dictionary named `loaders_scratch`. Your train data loader should be at `loaders_scratch['train']`, your validation data loader should be at `loaders_scratch['valid']`, and your test data loader should be at `loaders_scratch['test']`.

You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [ ]: # import stuff  
  
%matplotlib inline  
%config InlineBackend.figure_format = 'retina'  
  
import os  
  
import matplotlib.pyplot as plt  
  
import torch  
nn = torch.nn  
F = torch.nn.functional  
print('pytorch version: ',torch.__version__)  
  
from torchvision import datasets, transforms  
  
import numpy as np
```

```
np.random.seed(13) #static random numbers

import cv2

import splitfolders

import inspect
from datetime import datetime
```

pytorch version: 1.9.0+cu111

```
In [ ]: torch.cuda.empty_cache()
```

```
In [ ]:
try:
    import playsound
except:
    print('no playsound module')

def training_done():
    try:
        playsound.playsound('Krypton.mp3')
    except:
        print('no sound')
```

```
In [ ]:
data_train = 'landmark_images/train' #this is where the raw data is

dir_train = 'images/train'
dir_test = 'images/test'
dir_val= 'images/val'

if os.path.exists(dir_train) == False:
    splitfolders.ratio(data_train, output="images", seed=69, ratio=(.8, .1, .1), group_
```

```
In [ ]:
### TODO: Write data Loaders for training, validation, and test sets
## Specify appropriate transforms, and batch_sizes

# mean = [0.485, 0.456, 0.406]
# std = [0.229, 0.224, 0.225]

size = 128
mean = np.array([0.485, 0.456, 0.406])
std = np.array([0.229, 0.224, 0.225])

#https://pytorch.org/vision/stable/transforms.html
# (0.485, 0.456, 0.406), (0.229, 0.224, 0.225)

# train and valid
# maybe add a new more transforms
# https://pytorch.org/vision/stable/transforms.html
# transform0 = transforms.Compose([
#     transforms.RandomRotation(10),
#     transforms.Resize(size=(512,512)),
#     transforms.RandomResizedCrop(size=(256,256)),
#     # transforms.Resize(size=(128,128)),
```

```

#     # transforms.Resize(size=(64,64)),
#     # transforms.Grayscale(),
#     # transforms.RandomAdjustSharpness(1, p=1),
#     # transforms.RandomVerticalFlip(),
#     transforms.RandomHorizontalFlip(p=0.5),
#     transforms.ToTensor(),
#     # transforms.Normalize([0.5]*3,[0.5]*3),
#     ])

# # test
# transform1 = transforms.Compose([
#     # transforms.Resize(300),
#     transforms.Resize(size=(256,256)),
#     # transforms.Resize(size=(128,128)),
#     # transforms.Resize(size=(64,64)),
#     # transforms.Grayscale(),
#     transforms.ToTensor(),
#     # transforms.Normalize([0.5]*3,[0.5]*3),
#     ])

transform0 = transforms.Compose([
    transforms.RandomRotation(10),
    transforms.RandomResizedCrop(size),
    transforms.RandomHorizontalFlip(p=0.1),
    transforms.ColorJitter(brightness=0.05, contrast=0.05, saturation=0.5, hue=0.05),
    # transforms.ColorJitter(contrast=0.05),
    # transforms.ColorJitter(saturation=0.05),
    # transforms.ColorJitter(hue=0.05),
    transforms.ToTensor(),
    transforms.Normalize(mean, std)
])

transform1 = transforms.Compose([
    transforms.CenterCrop((size, size)),
    transforms.ToTensor(),
    transforms.Normalize(mean, std)
])

loaders_scratch = {}
loaders = [
    {
        'name':'train',
        'dir': dir_train,
        'transform': transform0,
    },
    {
        'name':'valid',
        'dir': dir_val,
        'transform': transform1,
    },
    {
        'name':'test',
        'dir': dir_test,
        'transform': transform1,
    },
]
for l in loaders:
    dataset = datasets.ImageFolder(root=l['dir'], transform=l['transform'])
    print(l['name'],len(dataset))

```

```
loaders_scratch[1['name']] = torch.utils.data.DataLoader(dataset, batch_size=64, sh

print(*loaders_scratch)

train 3996
valid 499
test 501
train valid test

In [ ]:
dataset = datasets.ImageFolder(root=dir_train, transform=transform0)
# print('\n'.join(dataset.classes))

label_map = list(dataset.classes)
print(*label_map, sep='\n')

00.Haleakala_National_Park
01.Mount_Rainier_National_Park
02.Ljubljana_Castle
03.Dead_Sea
04.Wroclaws_Dwarves
05.London_Olympic_Stadium
06.Niagara_Falls
07.Stonehenge
08.Grand_Canyon
09.Golden_Gate_Bridge
10.Edinburgh_Castle
11.Mount_Rushmore_National_Memorial
12.Kantanagar_Temple
13.Yellowstone_National_Park
14.Terminal_Tower
15.Central_Park
16.Eiffel_Tower
17.Changdeokgung
18.Delicate_Arch
19.Vienna_City_Hall
20.Matterhorn
21.Taj_Mahal
22.Moscow_Raceway
23.Externsteine
24.Soreq_Cave
25.Banff_National_Park
26.Pont_du_Gard
27.Seattle_Japanese_Garden
28.Sydney_Harbour_Bridge
29.Petronas_Towers
30.Brooklyn_Bridge
31.Washington_Monument
32.Hanging_Temple
33.Sydney_Opera_House
34.Great_Barrier_Reef
35.Monumento_a_la_Revolucion
36.Badlands_National_Park
37.Atomium
38.Forth_Bridge
39.Gateway_of_India
40.Stockholm_City_Hall
41.Machu_Picchu
42.Death_Valley_National_Park
43.Gullfoss_Falls
```

```
44.Trevi_Fountain  
45.Temple_of_Heaven  
46.Great_Wall_of_China  
47.Prague_Astronomical_Clock  
48.Whitby_Abbey  
49.Temple_of_Olympian_Zeus
```

resource for rendering image with plt

https://github.com/udacity/deep-learning-v2-pytorch/blob/master/convolutional-neural-networks/cifar-cnn/cifar10_cnn_solution.ipynb

resource for splitting the training and validation

<https://stackoverflow.com/questions/53074712/how-to-split-folder-of-images-into-test-training-validation-sets-with-stratified>

In []:

```
# helper function from  
# https://github.com/udacity/deep-Learning-v2-pytorch/blob/master/convolutional-neural-  
  
# helper function to un-normalize and display an image  
def imshow(image):  
    # img = img / 2 #+ 0.5 # unnormalize ..don't need this line  
    image = image.numpy() * std[:, None, None] + mean[:, None, None]  
    plt.imshow(np.transpose(image, (1, 2, 0))) # convert from Tensor image  
  
def show_images(images, labels, title=None, figsize=(20,5), image_limit = 10, color_map=None)  
  
    # unnormalize image  
    images = images.numpy() * std[:, None, None] + mean[:, None, None]  
    fig = plt.figure(figsize=figsize, facecolor='white')  
    fig.suptitle(title, fontsize=16)  
  
    rows = int(np.ceil(image_limit/5))  
  
    for idx in range(image_limit):  
        ax = fig.add_subplot(rows, 5, idx+1, xticks=[], yticks=[])  
        # imshow(images[idx])  
  
        plt.imshow(np.transpose(images[idx], (1, 2, 0)), cmap=color_map)  
        ax.set_title(label_map[labels[idx].item()], fontsize = 12)
```

Question 1: Describe your chosen procedure for preprocessing the data.

- How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why?
- Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

Answer:

- How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why?
 - i'm using the transforms within the torchvision library
 - i'm using
 - random rotation -10 and 10
 - random resize and crop (256 size)
 - random horizontal flip ~~* grey scaling them.. i don't think color will matter that much~~
 - random color jitter (brightness, contrast, etc)
 - and normalize ~~* random vertical flip~~
 - i'm cropping to a 256 square, because it's a square and divisible by 2.
 - we need to make sure all images are the same size for our NN, since our inputs have to be the same length.
- Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?
 - yes, we are rotating a bit, resizing cropping and flipping... this makes it so our AI doesn't get too used to the same exact image.

(IMPLEMENTATION) Visualize a Batch of Training Data

Use the code cell below to retrieve a batch of images from your train data loader, display at least 5 images simultaneously, and label each displayed image with its class name (e.g., "Golden Gate Bridge").

Visualizing the output of your data loader is a great way to ensure that your data loading and preprocessing are working as expected.

In []:

```
import matplotlib.pyplot as plt
%matplotlib inline

## TODO: visualize a batch of the train data Loader

## the class names can be accessed at the `classes` attribute
## of your dataset object (e.g., `train_dataset.classes`)

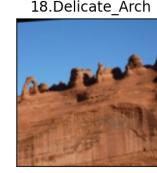
# train data

images, labels = next(iter(loaders_scratch['train']))
show_images(images, labels, title='Train', color_map=None)

# images, labels = next(iter(loaders_scratch['test']))
# show_images(images, labels, title='Test')

## used for random...but we are already shuffling
# r = np.random.choice(range(images.shape[0]))
# print(r)
# print(images[r].shape)
# imshow(images[r])
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



Initialize use_cuda variable

Steps

1. download & install the cuda toolkit
 - <https://developer.nvidia.com/cuda-toolkit-archive>
 - most newer versions should be fine
2. get the specific version of pytorch
 - <https://pytorch.org/get-started/locally/>
 - pick your OS, Cuda version, and etc
 - run the provided command

notes:

Cuda is only for specific Nvidia GPUs

In []:

```
# useful variable that tells us whether we should use the GPU
use_cuda = torch.cuda.is_available()
print(use_cuda)
```

True

(IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and fill in the function `get_optimizer_scratch` below.

In []:

```
## TODO: select loss function
criterion_scratch = None

# i do not have inputs of targets, yet
```

```

# so i'll use this function when the time comes
# def get_criterion_scratch(input,target):
#     criterion_scratch = F.nll_loss(input=input,target=target)

# we used this one in the lesson
criterion_scratch = nn.CrossEntropyLoss()

#others to try
# criterion_scratch = nn.L1Loss()
# criterion_scratch = nn.NLLLoss()
# criterion_scratch = nn.MSELoss()
# criterion_scratch = nn.HingeEmbeddingLoss()

# def get_optimizer_scratch(model, Learn_rate=0.001):
def get_optimizer_scratch(model):
    ## TODO: select and return an optimizer
    # optimizer = torch.optim.SGD(model.parameters(), lr = 0.001)

    # others to try
    # optimizer = torch.optim.RMSprop(model.parameters(), lr=0.001)
    # optimizer = torch.optim.Adadelta(model.parameters(), lr=0.001)
    optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
    # optimizer = torch.optim.Adagrad(model.parameters(), lr=0.001)

    return optimizer

```

(IMPLEMENTATION) Model Architecture

Create a CNN to classify images of landmarks. Use the template in the code cell below.

In []:

```

import torch.nn as nn

# define the CNN architecture
class Net(nn.Module):
    ## TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()

        ## Define layers of a CNN

        # for between
        self.pool = nn.MaxPool2d(2,2)

        # # input is 3 (RGB)
        # # output channels to 9 (how thick out tensor will be)
        # # kernel is (3,3) ...so we will be looking at a 3x3 pixels at a time
        # self.conv1 = nn.Conv2d(3,64,kernel_size=(3,3),stride=1,padding=1)
        # self.conv2 = nn.Conv2d(64,64,kernel_size=(3,3),stride=1,padding=1)
        # self.conv3 = nn.Conv2d(64,128,kernel_size=(3,3),stride=1,padding=1)
        # self.conv4 = nn.Conv2d(128,128,kernel_size=(3,3),stride=1,padding=1)
        # self.conv5 = nn.Conv2d(128,256,kernel_size=(3,3),stride=1,padding=1)
        # self.conv6 = nn.Conv2d(256,256,kernel_size=(3,3),stride=1,padding=1)
        # self.conv7 = nn.Conv2d(256,512,kernel_size=(3,3),stride=1,padding=1)
        # self.conv8 = nn.Conv2d(512,512,kernel_size=(3,3),stride=1,padding=1)

        self.conv1 = nn.Conv2d(3,16,kernel_size=(3,3),stride=1,padding=1)

```

```

self.conv2 = nn.Conv2d(16,32,kernel_size=(3,3),stride=1,padding=1)
self.conv3 = nn.Conv2d(32,64,kernel_size=(3,3),stride=1,padding=1)
# self.conv4 = nn.Conv2d(64,128,kernel_size=(3,3),stride=1,padding=1)

# self.fc1 = nn.Linear(131072,65536)
# self.fc2 = nn.Linear(65536,50)

self.fc1 = nn.Linear(16384,256)
self.fc2 = nn.Linear(256,128)
self.fc3 = nn.Linear(128,50)

self.dropout = nn.Dropout(0.2)

self.printed_x = False

def print_forward(self):
    return str(inspect.getsource(self.forward))

def print_model(self):
    return str(self)

def forward(self, x):
    ## Define forward behavior

    x = self.conv1(x)
    x = F.leaky_relu(x)
    x = self.pool(x)

    x = self.conv2(x)
    x = F.leaky_relu(x)
    x = self.pool(x)

    x = self.conv3(x)
    x = F.leaky_relu(x)
    x = self.pool(x)

    # x = self.conv4(x)
    # x = self.pool(x)
    # x = F.Leaky_relu(x)

    x = x.view(x.size(0),-1)

    if self.printed_x == False:
        print('x.shape: ',x.shape)
        self.printed_x = True
    #needed to find the correct shape for Line ^24

    x = self.dropout(x)
    x = F.leaky_relu(self.fc1(x),negative_slope=0.2)
    x = self.dropout(x)
    x = F.leaky_relu(self.fc2(x),negative_slope=0.2)
    x = self.dropout(x)
    x = F.leaky_relu(self.fc3(x),negative_slope=0.2)

    # x = F.relu(self.fc2(x))
    x = F.log_softmax(x,dim=1)

return x

```

```

#--# Do NOT modify the code below this line. #--#
# instantiate the CNN
model_scratch = Net()

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

```

Question 2: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer:

my first few models just doubled the convolutional layer's output and then had one fully connected layer, however this did not give me over 20%.

after looking at vgg and resnet i tried to mimic them. if these models had good scores and i should be able to mimic them and get a better output myself. so i added more fully connected layers, pools and dropouts.

after from trial and error i got to a 28% accuracy on the test function.

(IMPLEMENTATION) Implement the Training Algorithm

Implement your training algorithm in the code cell below. [Save the final model parameters](#) at the filepath stored in the variable `save_path`.

```
In [ ]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf

    for epoch in range(1, n_epochs+1):
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0

        #####
        # train the model #
        #####
        # set the module to training mode
        model.train()
        for batch_idx, (data, target) in enumerate(loaders['train']):
            # move to GPU

            if use_cuda:
                torch.cuda.empty_cache()
                data, target = data.cuda(), target.cuda()

            ## TODO: find the loss and update the model parameters accordingly
            ## record the average training loss, using something like
            ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data.item()) - t

            # this resets before back propagation
            optimizer.zero_grad()
```

```

# get model output
output = model(data)

# calculate loss
loss = criterion(output, target)
loss.backward()
optimizer.step()

# train based on loss
train_loss += ((1 / (batch_idx + 1)) * (loss.data.item() - train_loss))

#####
# validate the model #
#####
# set the model to evaluation mode
model.eval()
for batch_idx, (data, target) in enumerate(loaders['valid']):
    # move to GPU
    if use_cuda:
        torch.cuda.empty_cache()
        data, target = data.cuda(), target.cuda()

    ## TODO: update average validation loss
    output = model(data)
    # calculate the batch loss
    loss = criterion(output, target)
    # update average validation loss
    valid_loss += ((1 / (batch_idx + 1)) * (loss.data.item() - valid_loss))

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch,
    train_loss,
    valid_loss
))

## TODO: if the validation loss has decreased, save the model at the filepath s
if valid_loss <= valid_loss_min:

    # print('valid_loss: {:.10f}'.format(valid_loss))
    valid_loss_min = valid_loss
    print('valid_loss_min: {:.10f}'.format(valid_loss_min))

    # print('saving model')

    # save model
    torch.save(model.state_dict(), save_path)
    print('saved: ', save_path)

return model

```

(IMPLEMENTATION) Experiment with the Weight Initialization

Use the code cell below to define a custom weight initialization, and then train with your weight initialization for a few epochs. Make sure that neither the training loss nor validation loss is `nan`.

Later on, you will be able to see how this compares to training with PyTorch's default weight initialization.

```
In [ ]: def custom_weight_init(m):
    ## TODO: implement a weight initialization strategy

    #get the name of the layer
    name = m.__class__.__name__
    print(name)
    if name.find('Linear') != -1:

        # get the number of the inputs
        y = (1.0/np.sqrt(m.in_features))

        m.weight.data.normal_(0, y)
        m.bias.data.fill_(0)

    print(model_scratch)

#-#-# Do NOT modify the code below this Line. #-#-#

model_scratch.apply(custom_weight_init)
model_scratch = train(20, loaders_scratch, model_scratch, get_optimizer_scratch(model_s
    criterion_scratch, use_cuda, 'ignore.pt')
# model_scratch = train(1, Loaders_scratch, model_scratch, get_optimizer_scratch(model_
#     criterion_scratch, use_cuda, 'ignore.pt')

training_done()
```

```
Net(
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv1): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (fc1): Linear(in_features=16384, out_features=256, bias=True)
  (fc2): Linear(in_features=256, out_features=128, bias=True)
  (fc3): Linear(in_features=128, out_features=50, bias=True)
  (dropout): Dropout(p=0.2, inplace=False)
)
MaxPool2d
Conv2d
Conv2d
Conv2d
Linear
Linear
Linear
Dropout
Net
Epoch: 1      Training Loss: 3.547239      Validation Loss: 3.537234
valid_loss_min: 3.5372340977
saved: ignore.pt
Epoch: 2      Training Loss: 2.995708      Validation Loss: 3.318236
valid_loss_min: 3.3182357848
saved: ignore.pt
Epoch: 3      Training Loss: 2.753861      Validation Loss: 3.383212
Epoch: 4      Training Loss: 2.623263      Validation Loss: 3.209296
```

```

valid_loss_min: 3.2092956901
saved: ignore.pt
Epoch: 5      Training Loss: 2.536053      Validation Loss: 3.364543
Epoch: 6      Training Loss: 2.435655      Validation Loss: 3.158963
valid_loss_min: 3.1589625180
saved: ignore.pt
Epoch: 7      Training Loss: 2.359009      Validation Loss: 3.177190
Epoch: 8      Training Loss: 2.325564      Validation Loss: 3.142543
valid_loss_min: 3.1425430179
saved: ignore.pt
Epoch: 9      Training Loss: 2.230478      Validation Loss: 3.268974
Epoch: 10     Training Loss: 2.202910      Validation Loss: 3.285772
Epoch: 11     Training Loss: 2.175727      Validation Loss: 3.255863
Epoch: 12     Training Loss: 2.184670      Validation Loss: 3.155661
Epoch: 13     Training Loss: 2.128167      Validation Loss: 3.124331
valid_loss_min: 3.1243307889
saved: ignore.pt
Epoch: 14     Training Loss: 2.076077      Validation Loss: 3.252947
Epoch: 15     Training Loss: 2.073512      Validation Loss: 3.264821
Epoch: 16     Training Loss: 2.025115      Validation Loss: 3.266159
Epoch: 17     Training Loss: 2.006955      Validation Loss: 3.238283
Epoch: 18     Training Loss: 1.945982      Validation Loss: 3.219480
Epoch: 19     Training Loss: 1.948232      Validation Loss: 3.205019
Epoch: 20     Training Loss: 1.976405      Validation Loss: 3.189481

```

(IMPLEMENTATION) Train and Validate the Model

Run the next code cell to train your model.

```
In [ ]: #this allowed me to save information about my model

def save_model_specs(m):
    dttm = datetime.now().strftime("%Y%m%d%H%M%S")
    name = 'model_specs_' + dttm + '.txt'
    with open(name, 'w') as f:
        f.write(model_scratch.print_model())
        f.write('\n\n')
        f.write(model_scratch.print_forward())
    return name

# save_model_specs(model_scratch)
```

```
In [ ]: ## TODO: you may change the number of epochs if you'd like,
## but changing it is not required
num_epochs = 200

#save information about this model
ms = save_model_specs(model_scratch)

#---# Do NOT modify the code below this line. #---#

# function to re-initialize a model with pytorch's default weight initialization
def default_weight_init(m):
    reset_parameters = getattr(m, 'reset_parameters', None)
    if callable(reset_parameters):
        m.reset_parameters()
```

```

# reset the model parameters
model_scratch.apply(default_weight_init)

# train the model
model_scratch = train(num_epochs, loaders_scratch, model_scratch, get_optimizer_scratch
                      criterion_scratch, use_cuda, 'model_scratch.pt')

training_done()

```

C:\Users\JGarza\miniconda3\envs\lm3\lib\site-packages\torch\nn\functional.py:718: UserWarning: Named tensors and all their associated APIs are an experimental feature and subject to change. Please do not use them for anything important until they are released as stable. (Triggered internally at ..\c10\core/TensorImpl.h:1156.)

```

    return torch.max_pool2d(input, kernel_size, stride, padding, dilation, ceil_mode)
x.shape: torch.Size([64, 16384])
Epoch: 1      Training Loss: 3.873153          Validation Loss: 3.923265
valid_loss_min: 3.9232650101
saved: model_scratch.pt
Epoch: 2      Training Loss: 3.742517          Validation Loss: 3.779026
valid_loss_min: 3.7790262401
saved: model_scratch.pt
Epoch: 3      Training Loss: 3.610291          Validation Loss: 3.705194
valid_loss_min: 3.7051938176
saved: model_scratch.pt
Epoch: 4      Training Loss: 3.483709          Validation Loss: 3.653688
valid_loss_min: 3.6536877751
saved: model_scratch.pt
Epoch: 5      Training Loss: 3.383897          Validation Loss: 3.612159
valid_loss_min: 3.6121587157
saved: model_scratch.pt
Epoch: 6      Training Loss: 3.297176          Validation Loss: 3.596963
valid_loss_min: 3.5969630182
saved: model_scratch.pt
Epoch: 7      Training Loss: 3.195979          Validation Loss: 3.573895
valid_loss_min: 3.5738947093
saved: model_scratch.pt
Epoch: 8      Training Loss: 3.142290          Validation Loss: 3.531185
valid_loss_min: 3.5311848819
saved: model_scratch.pt
Epoch: 9      Training Loss: 3.080629          Validation Loss: 3.584989
Epoch: 10     Training Loss: 3.001242          Validation Loss: 3.446130
valid_loss_min: 3.4461302757
saved: model_scratch.pt
Epoch: 11     Training Loss: 2.946821          Validation Loss: 3.443898
valid_loss_min: 3.4438978732
saved: model_scratch.pt
Epoch: 12     Training Loss: 2.871822          Validation Loss: 3.474834
Epoch: 13     Training Loss: 2.848054          Validation Loss: 3.517810
Epoch: 14     Training Loss: 2.815887          Validation Loss: 3.506477
Epoch: 15     Training Loss: 2.761858          Validation Loss: 3.399811
valid_loss_min: 3.3998109698
saved: model_scratch.pt
Epoch: 16     Training Loss: 2.735817          Validation Loss: 3.385872
valid_loss_min: 3.3858715892
saved: model_scratch.pt
Epoch: 17     Training Loss: 2.676215          Validation Loss: 3.453596
Epoch: 18     Training Loss: 2.618847          Validation Loss: 3.448229
Epoch: 19     Training Loss: 2.593927          Validation Loss: 3.383564
valid_loss_min: 3.3835643530
saved: model_scratch.pt

```

```
Epoch: 20      Training Loss: 2.579453      Validation Loss: 3.354790
valid_loss_min: 3.3547903299
saved: model_scratch.pt
Epoch: 21      Training Loss: 2.563262      Validation Loss: 3.322260
valid_loss_min: 3.3222599030
saved: model_scratch.pt
Epoch: 22      Training Loss: 2.531193      Validation Loss: 3.354297
Epoch: 23      Training Loss: 2.506894      Validation Loss: 3.433205
Epoch: 24      Training Loss: 2.436046      Validation Loss: 3.337532
Epoch: 25      Training Loss: 2.461715      Validation Loss: 3.268325
valid_loss_min: 3.2683249116
saved: model_scratch.pt
Epoch: 26      Training Loss: 2.396691      Validation Loss: 3.381772
Epoch: 27      Training Loss: 2.396133      Validation Loss: 3.382679
Epoch: 28      Training Loss: 2.388649      Validation Loss: 3.332398
Epoch: 29      Training Loss: 2.340199      Validation Loss: 3.232119
valid_loss_min: 3.2321191430
saved: model_scratch.pt
Epoch: 30      Training Loss: 2.304878      Validation Loss: 3.313002
Epoch: 31      Training Loss: 2.296762      Validation Loss: 3.344792
Epoch: 32      Training Loss: 2.300024      Validation Loss: 3.358517
Epoch: 33      Training Loss: 2.259769      Validation Loss: 3.370558
Epoch: 34      Training Loss: 2.245814      Validation Loss: 3.281126
Epoch: 35      Training Loss: 2.210682      Validation Loss: 3.246249
Epoch: 36      Training Loss: 2.175223      Validation Loss: 3.217550
valid_loss_min: 3.2175504863
saved: model_scratch.pt
Epoch: 37      Training Loss: 2.171326      Validation Loss: 3.314161
Epoch: 38      Training Loss: 2.172530      Validation Loss: 3.329392
Epoch: 39      Training Loss: 2.173602      Validation Loss: 3.257709
Epoch: 40      Training Loss: 2.128484      Validation Loss: 3.305284
Epoch: 41      Training Loss: 2.155317      Validation Loss: 3.215880
valid_loss_min: 3.2158801556
saved: model_scratch.pt
Epoch: 42      Training Loss: 2.100343      Validation Loss: 3.345874
Epoch: 43      Training Loss: 2.029278      Validation Loss: 3.320589
Epoch: 44      Training Loss: 2.075221      Validation Loss: 3.314872
Epoch: 45      Training Loss: 2.063778      Validation Loss: 3.245261
Epoch: 46      Training Loss: 2.041589      Validation Loss: 3.247769
Epoch: 47      Training Loss: 2.045937      Validation Loss: 3.357670
Epoch: 48      Training Loss: 2.036675      Validation Loss: 3.336010
Epoch: 49      Training Loss: 1.999043      Validation Loss: 3.264808
Epoch: 50      Training Loss: 1.984028      Validation Loss: 3.300476
Epoch: 51      Training Loss: 1.943146      Validation Loss: 3.340531
Epoch: 52      Training Loss: 1.976503      Validation Loss: 3.290643
Epoch: 53      Training Loss: 1.994231      Validation Loss: 3.235343
Epoch: 54      Training Loss: 1.939168      Validation Loss: 3.267083
Epoch: 55      Training Loss: 1.916948      Validation Loss: 3.239802
Epoch: 56      Training Loss: 1.989915      Validation Loss: 3.222888
Epoch: 57      Training Loss: 1.932901      Validation Loss: 3.275119
Epoch: 58      Training Loss: 1.916238      Validation Loss: 3.323514
Epoch: 59      Training Loss: 1.940147      Validation Loss: 3.313124
Epoch: 60      Training Loss: 1.907206      Validation Loss: 3.290156
Epoch: 61      Training Loss: 1.942205      Validation Loss: 3.293176
Epoch: 62      Training Loss: 1.903101      Validation Loss: 3.185369
valid_loss_min: 3.1853693724
saved: model_scratch.pt
Epoch: 63      Training Loss: 1.884316      Validation Loss: 3.267317
Epoch: 64      Training Loss: 1.833592      Validation Loss: 3.346903
Epoch: 65      Training Loss: 1.840247      Validation Loss: 3.452234
Epoch: 66      Training Loss: 1.847975      Validation Loss: 3.248863
```

Epoch: 67	Training Loss: 1.835166	Validation Loss: 3.438620
Epoch: 68	Training Loss: 1.821640	Validation Loss: 3.320481
Epoch: 69	Training Loss: 1.837836	Validation Loss: 3.206787
Epoch: 70	Training Loss: 1.803670	Validation Loss: 3.399221
Epoch: 71	Training Loss: 1.831343	Validation Loss: 3.320581
Epoch: 72	Training Loss: 1.855379	Validation Loss: 3.279020
Epoch: 73	Training Loss: 1.785960	Validation Loss: 3.265032
Epoch: 74	Training Loss: 1.767898	Validation Loss: 3.314181
Epoch: 75	Training Loss: 1.786411	Validation Loss: 3.370351
Epoch: 76	Training Loss: 1.802403	Validation Loss: 3.350020
Epoch: 77	Training Loss: 1.725970	Validation Loss: 3.279673
Epoch: 78	Training Loss: 1.757657	Validation Loss: 3.254071
Epoch: 79	Training Loss: 1.721502	Validation Loss: 3.215011
Epoch: 80	Training Loss: 1.713209	Validation Loss: 3.346990
Epoch: 81	Training Loss: 1.701474	Validation Loss: 3.351555
Epoch: 82	Training Loss: 1.697753	Validation Loss: 3.249437
Epoch: 83	Training Loss: 1.764958	Validation Loss: 3.219951
Epoch: 84	Training Loss: 1.717125	Validation Loss: 3.282340
Epoch: 85	Training Loss: 1.679295	Validation Loss: 3.331624
Epoch: 86	Training Loss: 1.686862	Validation Loss: 3.340589
Epoch: 87	Training Loss: 1.658542	Validation Loss: 3.330867
Epoch: 88	Training Loss: 1.716238	Validation Loss: 3.301521
Epoch: 89	Training Loss: 1.655747	Validation Loss: 3.258401
Epoch: 90	Training Loss: 1.711830	Validation Loss: 3.256185
Epoch: 91	Training Loss: 1.698082	Validation Loss: 3.204439
Epoch: 92	Training Loss: 1.670694	Validation Loss: 3.312573
Epoch: 93	Training Loss: 1.686687	Validation Loss: 3.341612
Epoch: 94	Training Loss: 1.650589	Validation Loss: 3.136930
valid_loss_min:	3.1369304061	
saved:	model_scratch.pt	
Epoch: 95	Training Loss: 1.644793	Validation Loss: 3.304775
Epoch: 96	Training Loss: 1.668745	Validation Loss: 3.396186
Epoch: 97	Training Loss: 1.619471	Validation Loss: 3.188164
Epoch: 98	Training Loss: 1.626114	Validation Loss: 3.251437
Epoch: 99	Training Loss: 1.594419	Validation Loss: 3.270999
Epoch: 100	Training Loss: 1.604230	Validation Loss: 3.199820
Epoch: 101	Training Loss: 1.625020	Validation Loss: 3.243446
Epoch: 102	Training Loss: 1.637467	Validation Loss: 3.493325
Epoch: 103	Training Loss: 1.624048	Validation Loss: 3.343145
Epoch: 104	Training Loss: 1.580080	Validation Loss: 3.235943
Epoch: 105	Training Loss: 1.626861	Validation Loss: 3.318541
Epoch: 106	Training Loss: 1.608351	Validation Loss: 3.351922
Epoch: 107	Training Loss: 1.627345	Validation Loss: 3.319736
Epoch: 108	Training Loss: 1.618487	Validation Loss: 3.294760
Epoch: 109	Training Loss: 1.581394	Validation Loss: 3.363090
Epoch: 110	Training Loss: 1.582254	Validation Loss: 3.323109
Epoch: 111	Training Loss: 1.553928	Validation Loss: 3.387664
Epoch: 112	Training Loss: 1.578295	Validation Loss: 3.402485
Epoch: 113	Training Loss: 1.562038	Validation Loss: 3.319520
Epoch: 114	Training Loss: 1.602672	Validation Loss: 3.337283
Epoch: 115	Training Loss: 1.566514	Validation Loss: 3.341435
Epoch: 116	Training Loss: 1.574522	Validation Loss: 3.320448
Epoch: 117	Training Loss: 1.544316	Validation Loss: 3.348370
Epoch: 118	Training Loss: 1.553516	Validation Loss: 3.383344
Epoch: 119	Training Loss: 1.550338	Validation Loss: 3.277851
Epoch: 120	Training Loss: 1.536621	Validation Loss: 3.298728
Epoch: 121	Training Loss: 1.517157	Validation Loss: 3.275847
Epoch: 122	Training Loss: 1.523779	Validation Loss: 3.397568
Epoch: 123	Training Loss: 1.510676	Validation Loss: 3.460572
Epoch: 124	Training Loss: 1.530161	Validation Loss: 3.241450
Epoch: 125	Training Loss: 1.502001	Validation Loss: 3.590507

Epoch: 126	Training Loss: 1.507606	Validation Loss: 3.355292
Epoch: 127	Training Loss: 1.504227	Validation Loss: 3.275490
Epoch: 128	Training Loss: 1.467714	Validation Loss: 3.334708
Epoch: 129	Training Loss: 1.481764	Validation Loss: 3.398428
Epoch: 130	Training Loss: 1.526085	Validation Loss: 3.467800
Epoch: 131	Training Loss: 1.494790	Validation Loss: 3.347629
Epoch: 132	Training Loss: 1.505155	Validation Loss: 3.201993
Epoch: 133	Training Loss: 1.531405	Validation Loss: 3.359160
Epoch: 134	Training Loss: 1.470475	Validation Loss: 3.418785
Epoch: 135	Training Loss: 1.460317	Validation Loss: 3.313516
Epoch: 136	Training Loss: 1.482611	Validation Loss: 3.276762
Epoch: 137	Training Loss: 1.484042	Validation Loss: 3.368085
Epoch: 138	Training Loss: 1.483218	Validation Loss: 3.506397
Epoch: 139	Training Loss: 1.468984	Validation Loss: 3.635416
Epoch: 140	Training Loss: 1.503921	Validation Loss: 3.301408
Epoch: 141	Training Loss: 1.453242	Validation Loss: 3.343073
Epoch: 142	Training Loss: 1.454843	Validation Loss: 3.313228
Epoch: 143	Training Loss: 1.470952	Validation Loss: 3.533153
Epoch: 144	Training Loss: 1.467885	Validation Loss: 3.395523
Epoch: 145	Training Loss: 1.443691	Validation Loss: 3.250160
Epoch: 146	Training Loss: 1.470001	Validation Loss: 3.413421
Epoch: 147	Training Loss: 1.379275	Validation Loss: 3.517614
Epoch: 148	Training Loss: 1.422496	Validation Loss: 3.417979
Epoch: 149	Training Loss: 1.434433	Validation Loss: 3.286021
Epoch: 150	Training Loss: 1.457590	Validation Loss: 3.449401
Epoch: 151	Training Loss: 1.465470	Validation Loss: 3.335582
Epoch: 152	Training Loss: 1.414177	Validation Loss: 3.283046
Epoch: 153	Training Loss: 1.406939	Validation Loss: 3.409252
Epoch: 154	Training Loss: 1.402871	Validation Loss: 3.460306
Epoch: 155	Training Loss: 1.437788	Validation Loss: 3.324561
Epoch: 156	Training Loss: 1.391447	Validation Loss: 3.299711
Epoch: 157	Training Loss: 1.465664	Validation Loss: 3.403409
Epoch: 158	Training Loss: 1.507976	Validation Loss: 3.264173
Epoch: 159	Training Loss: 1.384741	Validation Loss: 3.289305
Epoch: 160	Training Loss: 1.447772	Validation Loss: 3.285047
Epoch: 161	Training Loss: 1.358992	Validation Loss: 3.377420
Epoch: 162	Training Loss: 1.430509	Validation Loss: 3.429964
Epoch: 163	Training Loss: 1.375167	Validation Loss: 3.294994
Epoch: 164	Training Loss: 1.434617	Validation Loss: 3.304188
Epoch: 165	Training Loss: 1.427945	Validation Loss: 3.301825
Epoch: 166	Training Loss: 1.427787	Validation Loss: 3.387916
Epoch: 167	Training Loss: 1.378402	Validation Loss: 3.339757
Epoch: 168	Training Loss: 1.414967	Validation Loss: 3.387617
Epoch: 169	Training Loss: 1.397306	Validation Loss: 3.156665
Epoch: 170	Training Loss: 1.379892	Validation Loss: 3.386069
Epoch: 171	Training Loss: 1.421549	Validation Loss: 3.396428
Epoch: 172	Training Loss: 1.399397	Validation Loss: 3.203363
Epoch: 173	Training Loss: 1.409170	Validation Loss: 3.362925
Epoch: 174	Training Loss: 1.369631	Validation Loss: 3.303527
Epoch: 175	Training Loss: 1.353078	Validation Loss: 3.311549
Epoch: 176	Training Loss: 1.344728	Validation Loss: 3.378529
Epoch: 177	Training Loss: 1.388703	Validation Loss: 3.326973
Epoch: 178	Training Loss: 1.385569	Validation Loss: 3.391251
Epoch: 179	Training Loss: 1.378590	Validation Loss: 3.386482
Epoch: 180	Training Loss: 1.358225	Validation Loss: 3.192385
Epoch: 181	Training Loss: 1.374148	Validation Loss: 3.430165
Epoch: 182	Training Loss: 1.345461	Validation Loss: 3.312873
Epoch: 183	Training Loss: 1.392624	Validation Loss: 3.359835
Epoch: 184	Training Loss: 1.377776	Validation Loss: 3.293821
Epoch: 185	Training Loss: 1.403242	Validation Loss: 3.373683
Epoch: 186	Training Loss: 1.362348	Validation Loss: 3.293807

Epoch: 187	Training Loss: 1.370734	Validation Loss: 3.209247
Epoch: 188	Training Loss: 1.363375	Validation Loss: 3.318186
Epoch: 189	Training Loss: 1.332865	Validation Loss: 3.364515
Epoch: 190	Training Loss: 1.359444	Validation Loss: 3.462594
Epoch: 191	Training Loss: 1.326493	Validation Loss: 3.332458
Epoch: 192	Training Loss: 1.374573	Validation Loss: 3.553651
Epoch: 193	Training Loss: 1.374436	Validation Loss: 3.408561
Epoch: 194	Training Loss: 1.385739	Validation Loss: 3.306505
Epoch: 195	Training Loss: 1.375677	Validation Loss: 3.529334
Epoch: 196	Training Loss: 1.346959	Validation Loss: 3.308462
Epoch: 197	Training Loss: 1.323029	Validation Loss: 3.339844
Epoch: 198	Training Loss: 1.332034	Validation Loss: 3.374526
Epoch: 199	Training Loss: 1.318532	Validation Loss: 3.416389
Epoch: 200	Training Loss: 1.382229	Validation Loss: 3.215977

(IMPLEMENTATION) Test the Model

Run the code cell below to try out your model on the test dataset of landmark images. Run the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 20%.

```
In [ ]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    # set the module to evaluation mode
    model.eval()

    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            torch.cuda.empty_cache()
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data.item() - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %d%% (%d/%d)' % (
        100. * correct / total, correct, total))

    #if the model specs file is still there i can add the test scores to it
    try:
        with open(ms,'a') as f:
            f.write('\n\n')
            f.write('Test Loss: {:.6f}\n'.format(test_loss))
            f.write('\nTest Accuracy: %d%% (%d/%d)' % (
```

```

        100. * correct / total, correct, total))
except:
    print('no ms')
    pass

# Load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch.pt'))
test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)

```

Test Loss: 3.197581

Test Accuracy: 28% (144/501)
no ms

Step 2: Create a CNN to Classify Landmarks (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify landmarks from images. Your CNN must attain at least 60% accuracy on the test set.

(IMPLEMENTATION) Specify Data Loaders for the Landmark Dataset

Use the code cell below to create three separate [data loaders](#): one for training data, one for validation data, and one for test data. Randomly split the images located at `landmark_images/train` to create the train and validation data loaders, and use the images located at `landmark_images/test` to create the test data loader.

All three of your data loaders should be accessible via a dictionary named `loaders_transfer`. Your train data loader should be at `loaders_transfer['train']`, your validation data loader should be at `loaders_transfer['valid']`, and your test data loader should be at `loaders_transfer['test']`.

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

In []:

```

### TODO: Write data loaders for training, validation, and test sets
## Specify appropriate transforms, and batch_sizes
import os
import numpy as np
import torch
import torch.nn as nn

import torchvision
from torchvision import datasets, models, transforms
import matplotlib.pyplot as plt

%matplotlib inline

```

```
# Loaders_transfer = {'train': None, 'valid': None, 'test': None}
# Loaders_transfer = Loaders_scratch
```

```
In [ ]:
data_train = 'landmark_images/train' #this is where the raw data is

dir_train = 'images/train'
dir_test = 'images/test'
dir_val= 'images/val'

if os.path.exists(dir_train) == False:
    splitfolders.ratio(data_train, output="images", seed=69, ratio=(.8, .1, .1), group_
```

```
In [ ]:
size = 256
mean = np.array([0.485, 0.456, 0.406])
std = np.array([0.229, 0.224, 0.225])

transform0 = transforms.Compose([
    transforms.RandomRotation(30),
    transforms.RandomResizedCrop(size),
    transforms.RandomHorizontalFlip(p=0.1),
    transforms.ColorJitter(brightness=0.05, contrast=0.05, saturation=0.5, hue=0.05),
    # transforms.ColorJitter(contrast=0.05),
    # transforms.ColorJitter(saturation=0.05),
    # transforms.ColorJitter(hue=0.05),
    transforms.ToTensor(),
    transforms.Normalize(mean, std)
])

transform1 = transforms.Compose([
    transforms.CenterCrop((size, size)),
    transforms.ToTensor(),
    transforms.Normalize(mean, std)
])

loaders_transfer = {}
loaders = [
    {
        'name':'train',
        'dir': dir_train,
        'transform': transform0,
    },
    {
        'name':'valid',
        'dir': dir_val,
        'transform': transform1,
    },
    {
        'name':'test',
        'dir': dir_test,
        'transform': transform1,
    },
]

for l in loaders:
    dataset = datasets.ImageFolder(root=l['dir'], transform=l['transform'])
    print(l['name'], len(dataset))
```

```
loaders_transfer[1['name']] = torch.utils.data.DataLoader(dataset, batch_size=16, s  
  
print(*loaders_transfer)  
  
train 3996  
valid 499  
test 501  
train valid test
```

(IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and fill in the function `get_optimizer_transfer` below.

```
In [ ]:  
## TODO: select loss function  
# criterion_scratch = None  
  
# we used this one in the lesson  
criterion_transfer = nn.CrossEntropyLoss()  
  
#others to try  
# criterion_transfer = nn.L1Loss()  
# criterion_transfer = nn.NLLLoss()  
# criterion_transfer = nn.MSELoss()  
# criterion_transfer = nn.HingeEmbeddingLoss()  
  
# def get_optimizer_scratch(model, Learn_rate=0.001):  
def get_optimizer_transfer(model):  
    ## TODO: select and return an optimizer  
    optimizer = torch.optim.SGD(model.parameters(), lr = 0.001)  
  
    # others to try  
    # optimizer = torch.optim.RMSprop(model.parameters(), lr=0.001)  
    # optimizer = torch.optim.Adadelta(model.parameters(), lr=0.001)  
    # optimizer = torch.optim.Adam(model.parameters(), lr=0.001)  
    # optimizer = torch.optim.Adagrad(model.parameters(), lr=0.001)  
  
    return optimizer
```

(IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify images of landmarks. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```
In [ ]:  
# useful variable that tells us whether we should use the GPU  
use_cuda = torch.cuda.is_available()  
print(use_cuda)
```

True

```
In [ ]:  
torch.cuda.empty_cache()
```

```
In [ ]: ## TODO: Specify model architecture

# https://pytorch.org/vision/stable/models.html
# model_transfer = models.vgg16(pretrained=True) #not working only about 52%

model_transfer = models.vgg19_bn(pretrained=True)

# model_transfer = models.vgg11(pretrained=True)

for param in model_transfer.features.parameters():
    param.requires_grad = False

num_inputs = model_transfer.classifier[6].in_features
final_layer = nn.Linear(num_inputs, 50)
model_transfer.classifier[6] = final_layer

# model_transfer = models.resnet50(pretrained=True) #resnet is difficult to get working
# num_features = model_transfer.fc.in_features
# model_transfer.fc = nn.Linear(num_features, 50)

torch.save(model_transfer.state_dict(), 'model_transfer.pt')

#--# Do NOT modify the code below this Line. #-#-

if use_cuda:
    model_transfer = model_transfer.cuda()
```

Downloading: "https://download.pytorch.org/models/vgg19_bn-c79401a0.pth" to C:\Users\JGara/.cache\torch\hub\checkpoints\vgg19_bn-c79401a0.pth
100.0%

Question 3: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer:

~~we used the vgg16 in one of the demos, so i tried that one first and it seems to work ok.

after importing vgg16 we had to set the number of inputs based on the number of features, and then we set the outputs to 50 for our classes.~~

I tried a few models, and some were a bit difficult to set up. but after i saw the training on the vgg models i was able to get that model working.

vgg19 architecture is able to scale down the image in multiple layers...this allows it to retain the important parts of the image, and remove the noise. in addition the 3 fully connected layers are able to make more sense out of a larger number of inputs than 1 fully connected layer alone.

(IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath '`model_transfer.pt`' .

```
In [ ]: # TODO: train the model and save the best model parameters at filepath 'model_transfer.pt'
```

```
model_transfer = train(100, loaders_transfer, model_transfer, get_optimizer_transfer(model_transfer))

#---# Do NOT modify the code below this Line. #---#

# Load the model that got the best validation accuracy
model_transfer.load_state_dict(torch.load('model_transfer.pt'))
```

```
Epoch: 1      Training Loss: 3.801741      Validation Loss: 3.542806
valid_loss_min: 3.5428062156
saved: model_transfer.pt
Epoch: 2      Training Loss: 3.452338      Validation Loss: 3.166952
valid_loss_min: 3.1669523418
saved: model_transfer.pt
Epoch: 3      Training Loss: 3.122536      Validation Loss: 2.797232
valid_loss_min: 2.7972324118
saved: model_transfer.pt
Epoch: 4      Training Loss: 2.807328      Validation Loss: 2.502496
valid_loss_min: 2.5024958029
saved: model_transfer.pt
Epoch: 5      Training Loss: 2.514951      Validation Loss: 2.214901
valid_loss_min: 2.2149012610
saved: model_transfer.pt
Epoch: 6      Training Loss: 2.319168      Validation Loss: 2.071080
valid_loss_min: 2.0710796006
saved: model_transfer.pt
Epoch: 7      Training Loss: 2.146826      Validation Loss: 1.901805
valid_loss_min: 1.9018054828
saved: model_transfer.pt
Epoch: 8      Training Loss: 2.039973      Validation Loss: 1.811968
valid_loss_min: 1.8119676039
saved: model_transfer.pt
Epoch: 9      Training Loss: 1.941100      Validation Loss: 1.759248
valid_loss_min: 1.7592480369
saved: model_transfer.pt
Epoch: 10     Training Loss: 1.868742      Validation Loss: 1.674663
valid_loss_min: 1.6746630333
saved: model_transfer.pt
Epoch: 11     Training Loss: 1.834154      Validation Loss: 1.628738
valid_loss_min: 1.6287379581
saved: model_transfer.pt
Epoch: 12     Training Loss: 1.755940      Validation Loss: 1.598254
valid_loss_min: 1.5982540529
saved: model_transfer.pt
Epoch: 13     Training Loss: 1.745586      Validation Loss: 1.604606
Epoch: 14     Training Loss: 1.708091      Validation Loss: 1.532228
valid_loss_min: 1.5322275423
saved: model_transfer.pt
Epoch: 15     Training Loss: 1.661741      Validation Loss: 1.530211
valid_loss_min: 1.5302106105
saved: model_transfer.pt
Epoch: 16     Training Loss: 1.671059      Validation Loss: 1.499456
valid_loss_min: 1.4994556680
saved: model_transfer.pt
Epoch: 17     Training Loss: 1.602576      Validation Loss: 1.505682
Epoch: 18     Training Loss: 1.589733      Validation Loss: 1.481229
valid_loss_min: 1.4812291171
saved: model_transfer.pt
```

```
Epoch: 19      Training Loss: 1.567645      Validation Loss: 1.486745
Epoch: 20      Training Loss: 1.526702      Validation Loss: 1.447163
valid_loss_min: 1.4471627055
saved: model_transfer.pt
Epoch: 21      Training Loss: 1.530097      Validation Loss: 1.472251
Epoch: 22      Training Loss: 1.550866      Validation Loss: 1.424655
valid_loss_min: 1.4246554282
saved: model_transfer.pt
Epoch: 23      Training Loss: 1.524984      Validation Loss: 1.441221
Epoch: 24      Training Loss: 1.513371      Validation Loss: 1.396163
valid_loss_min: 1.3961628648
saved: model_transfer.pt
Epoch: 25      Training Loss: 1.467710      Validation Loss: 1.396367
Epoch: 26      Training Loss: 1.482894      Validation Loss: 1.398675
Epoch: 27      Training Loss: 1.450101      Validation Loss: 1.400391
Epoch: 28      Training Loss: 1.465930      Validation Loss: 1.418110
Epoch: 29      Training Loss: 1.450469      Validation Loss: 1.411157
Epoch: 30      Training Loss: 1.402869      Validation Loss: 1.443262
Epoch: 31      Training Loss: 1.399429      Validation Loss: 1.370469
valid_loss_min: 1.3704686575
saved: model_transfer.pt
Epoch: 32      Training Loss: 1.371501      Validation Loss: 1.384919
Epoch: 33      Training Loss: 1.372575      Validation Loss: 1.364801
valid_loss_min: 1.3648011656
saved: model_transfer.pt
Epoch: 34      Training Loss: 1.368450      Validation Loss: 1.390200
Epoch: 35      Training Loss: 1.349220      Validation Loss: 1.377157
Epoch: 36      Training Loss: 1.377531      Validation Loss: 1.415021
Epoch: 37      Training Loss: 1.360668      Validation Loss: 1.382934
Epoch: 38      Training Loss: 1.344022      Validation Loss: 1.354084
valid_loss_min: 1.3540841844
saved: model_transfer.pt
Epoch: 39      Training Loss: 1.354069      Validation Loss: 1.420224
Epoch: 40      Training Loss: 1.344322      Validation Loss: 1.352469
valid_loss_min: 1.3524686769
saved: model_transfer.pt
Epoch: 41      Training Loss: 1.309006      Validation Loss: 1.371570
Epoch: 42      Training Loss: 1.310723      Validation Loss: 1.356746
Epoch: 43      Training Loss: 1.306272      Validation Loss: 1.415458
Epoch: 44      Training Loss: 1.308766      Validation Loss: 1.363453
Epoch: 45      Training Loss: 1.284993      Validation Loss: 1.358798
Epoch: 46      Training Loss: 1.316673      Validation Loss: 1.362160
Epoch: 47      Training Loss: 1.305400      Validation Loss: 1.363035
Epoch: 48      Training Loss: 1.274133      Validation Loss: 1.386637
Epoch: 49      Training Loss: 1.284394      Validation Loss: 1.387577
Epoch: 50      Training Loss: 1.261272      Validation Loss: 1.352565
Epoch: 51      Training Loss: 1.292438      Validation Loss: 1.370316
Epoch: 52      Training Loss: 1.233448      Validation Loss: 1.305114
valid_loss_min: 1.3051141967
saved: model_transfer.pt
Epoch: 53      Training Loss: 1.247346      Validation Loss: 1.343543
Epoch: 54      Training Loss: 1.245966      Validation Loss: 1.326711
Epoch: 55      Training Loss: 1.258307      Validation Loss: 1.291160
valid_loss_min: 1.2911600727
saved: model_transfer.pt
Epoch: 56      Training Loss: 1.242302      Validation Loss: 1.308448
Epoch: 57      Training Loss: 1.253893      Validation Loss: 1.333049
Epoch: 58      Training Loss: 1.237516      Validation Loss: 1.301190
Epoch: 59      Training Loss: 1.225134      Validation Loss: 1.316601
Epoch: 60      Training Loss: 1.220613      Validation Loss: 1.343164
Epoch: 61      Training Loss: 1.210511      Validation Loss: 1.337858
```

```
Epoch: 62      Training Loss: 1.200453      Validation Loss: 1.310924
Epoch: 63      Training Loss: 1.207261      Validation Loss: 1.316450
Epoch: 64      Training Loss: 1.212559      Validation Loss: 1.352561
Epoch: 65      Training Loss: 1.187202      Validation Loss: 1.315956
Epoch: 66      Training Loss: 1.178044      Validation Loss: 1.316500
Epoch: 67      Training Loss: 1.201407      Validation Loss: 1.337132
Epoch: 68      Training Loss: 1.155121      Validation Loss: 1.338544
Epoch: 69      Training Loss: 1.193997      Validation Loss: 1.365828
Epoch: 70      Training Loss: 1.186053      Validation Loss: 1.286068
valid_loss_min: 1.2860682206
saved: model_transfer.pt
Epoch: 71      Training Loss: 1.167307      Validation Loss: 1.289281
Epoch: 72      Training Loss: 1.195489      Validation Loss: 1.272583
valid_loss_min: 1.2725828513
saved: model_transfer.pt
Epoch: 73      Training Loss: 1.167043      Validation Loss: 1.318778
Epoch: 74      Training Loss: 1.139039      Validation Loss: 1.339745
Epoch: 75      Training Loss: 1.142560      Validation Loss: 1.337402
Epoch: 76      Training Loss: 1.179030      Validation Loss: 1.290750
Epoch: 77      Training Loss: 1.136127      Validation Loss: 1.275186
Epoch: 78      Training Loss: 1.170927      Validation Loss: 1.277169
Epoch: 79      Training Loss: 1.117672      Validation Loss: 1.299860
Epoch: 80      Training Loss: 1.129766      Validation Loss: 1.320111
Epoch: 81      Training Loss: 1.134553      Validation Loss: 1.290523
Epoch: 82      Training Loss: 1.109167      Validation Loss: 1.309192
Epoch: 83      Training Loss: 1.158824      Validation Loss: 1.353063
Epoch: 84      Training Loss: 1.127334      Validation Loss: 1.359118
Epoch: 85      Training Loss: 1.111767      Validation Loss: 1.297525
Epoch: 86      Training Loss: 1.103297      Validation Loss: 1.294292
Epoch: 87      Training Loss: 1.123373      Validation Loss: 1.292377
Epoch: 88      Training Loss: 1.115293      Validation Loss: 1.304395
Epoch: 89      Training Loss: 1.096108      Validation Loss: 1.291939
Epoch: 90      Training Loss: 1.095606      Validation Loss: 1.293328
Epoch: 91      Training Loss: 1.079823      Validation Loss: 1.302960
Epoch: 92      Training Loss: 1.079580      Validation Loss: 1.302771
Epoch: 93      Training Loss: 1.088439      Validation Loss: 1.322767
Epoch: 94      Training Loss: 1.097192      Validation Loss: 1.290362
Epoch: 95      Training Loss: 1.119136      Validation Loss: 1.268183
valid_loss_min: 1.2681831997
saved: model_transfer.pt
Epoch: 96      Training Loss: 1.066409      Validation Loss: 1.262901
valid_loss_min: 1.2629005546
saved: model_transfer.pt
Epoch: 97      Training Loss: 1.075925      Validation Loss: 1.260939
valid_loss_min: 1.2609385019
saved: model_transfer.pt
Epoch: 98      Training Loss: 1.084537      Validation Loss: 1.270316
Epoch: 99      Training Loss: 1.072272      Validation Loss: 1.286101
Epoch: 100     Training Loss: 1.060633      Validation Loss: 1.265338
<All keys matched successfully>
```

Out[]:

(IMPLEMENTATION) Test the Model

Try out your model on the test dataset of landmark images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

In []:

```
test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

Test Loss: 1.343670

Test Accuracy: 65% (328/501)
no ms

Step 3: Write Your Landmark Prediction Algorithm

Great job creating your CNN models! Now that you have put in all the hard work of creating accurate classifiers, let's define some functions to make it easy for others to use your classifiers.

(IMPLEMENTATION) Write Your Algorithm, Part 1

Implement the function `predict_landmarks`, which accepts a file path to an image and an integer k, and then predicts the **top k most likely landmarks**. You are **required** to use your transfer learned CNN from Step 2 to predict the landmarks.

An example of the expected behavior of `predict_landmarks` :

```
>>> predicted_landmarks = predict_landmarks('example_image.jpg', 3)
>>> print(predicted_landmarks)
['Golden Gate Bridge', 'Brooklyn Bridge', 'Sydney Harbour Bridge']
```

```
In [ ]: import cv2
from PIL import Image

## the class names can be accessed at the `classes` attribute
## of your dataset object (e.g., `train_dataset.classes`)

def predict_landmarks(img_path, k):
    ## TODO: return the names of the top k Landmarks predicted by the transfer Learned

    image = Image.open(img_path)
    i = transform0(image)
    # imshow(i)
    i = i[None,:]
    if use_cuda:
        i = i.cuda()
    p = model_transfer(i).cpu()
    p = p.detach().numpy()[0]

    p_sorted = p.copy()
    p_sorted.sort()
    p_sorted = p_sorted[-k:]
    p_sorted = np.flip(p_sorted)

    # print(p)
    # print(p_sorted)

    targets = list(range(k))
    for index,i in enumerate(p):
        if i in p_sorted:
            pos = np.where(p_sorted == i)[0]
            # print(pos)
```

```

        targets[pos[0]] = {'position':pos[0], 'index':index, 'value':i, 'label':label}
    # print(*targets,sep='\n')
    return targets

# test on a sample image
predict_landmarks('images/test/09.Golden_Gate_Bridge/06dcc95353f874d9.jpg', 5)

```

```

Out[ ]: [{"position": 0,
           "index": 9,
           "value": 14.619126,
           "label": '09.Golden_Gate_Bridge'},
          {"position": 1,
           "index": 30,
           "value": 9.5372715,
           "label": '30.Brooklyn_Bridge'},
          {"position": 2, "index": 38, "value": 4.770165, "label": '38.Forth_Bridge'},
          {"position": 3,
           "index": 28,
           "value": 4.251987,
           "label": '28.Sydney_Harbour_Bridge'},
          {"position": 4,
           "index": 35,
           "value": 3.1836703,
           "label": '35.Monumento_a_la_Revolucion'}]

```

(IMPLEMENTATION) Write Your Algorithm, Part 2

In the code cell below, implement the function `suggest_locations`, which accepts a file path to an image as input, and then displays the image and the **top 3 most likely landmarks** as predicted by `predict_landmarks`.

Some sample output for `suggest_locations` is provided below, but feel free to design your own user experience!

```

In [ ]: def suggest_locations(img_path):
    # get Landmark predictions

    print('img_path: ',img_path)

    predicted_landmarks_list = predict_landmarks(img_path, 3)
    # print(type(predicted_Landmarks_list))

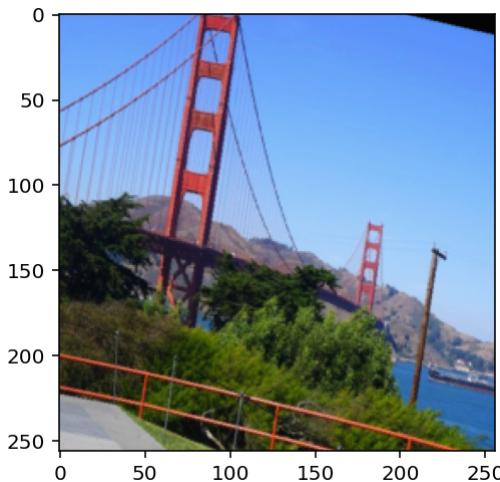
    ## TODO: display image and display Landmark predictions
    image = Image.open(img_path)
    i = transform0(image)
    imshow(i)

    print(*predicted_landmarks_list,sep='\n')

```

```
# test on a sample image
suggest_locations('images/test/09.Golden_Gate_Bridge/06dcc95353f874d9.jpg')
```

```
img_path: images/test/09.Golden_Gate_Bridge/06dcc95353f874d9.jpg
{'position': 0, 'index': 9, 'value': 16.038969, 'label': '09.Golden_Gate_Bridge'}
{'position': 1, 'index': 30, 'value': 9.797019, 'label': '30.Brooklyn_Bridge'}
{'position': 2, 'index': 38, 'value': 6.713119, 'label': '38.Forth_Bridge'}
```



(IMPLEMENTATION) Test Your Algorithm

Test your algorithm by running the `suggest_locations` function on at least four images on your computer. Feel free to use any images you like.

Question 4: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: (Three possible points for improvement)

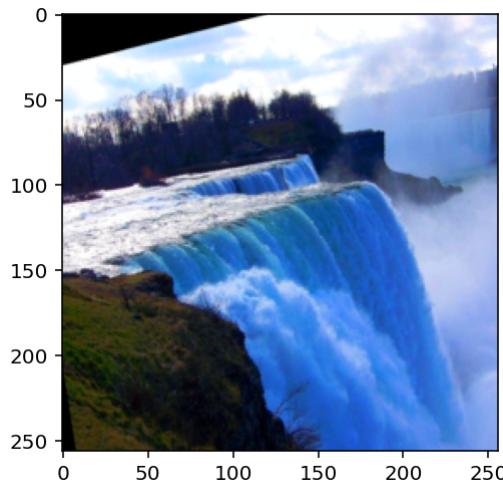
- i might be shrinking down the images too much on the transforms should it be 64? 128?
...more?
- my random adjustments might be too much
- i might not have enough convd layers
- i might need to have multiple fully Connected layers

```
In [ ]:
## TODO: Execute the `suggest_Locations` function on
## at least 4 images on your computer.
## Feel free to use as many code cells as needed.
```

```
suggest_locations('images/test/06.Niagara_Falls/35a6283e9c0a5d40.jpg')
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

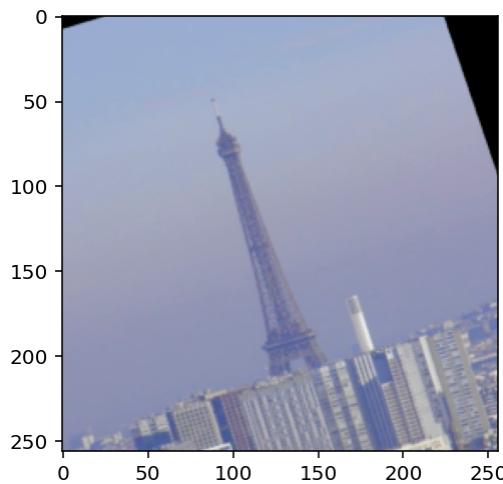
```
img_path: images/test/06.Niagara_Falls/35a6283e9c0a5d40.jpg
{'position': 0, 'index': 6, 'value': 10.27085, 'label': '06.Niagara_Falls'}
{'position': 1, 'index': 43, 'value': 9.590778, 'label': '43.Gullfoss_Falls'}
{'position': 2, 'index': 25, 'value': 5.131311, 'label': '25.Banff_National_Park'}
```



```
In [ ]: suggest_locations('images/test/16.Eiffel_Tower/5cd6b1c34778255c.jpg')
```

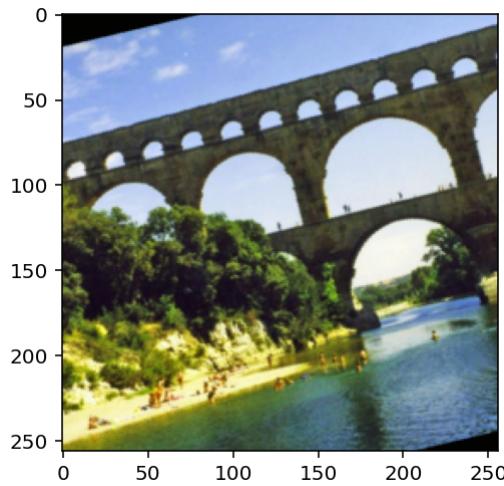
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```
img_path: images/test/16.Eiffel_Tower/5cd6b1c34778255c.jpg
{'position': 0, 'index': 31, 'value': 7.397161, 'label': '31.Washington_Monument'}
{'position': 1, 'index': 16, 'value': 6.9582233, 'label': '16.Eiffel_Tower'}
{'position': 2, 'index': 29, 'value': 6.919037, 'label': '29.Petronas_Towers'}
```



```
In [ ]: suggest_locations('images/test/26.Pont_du_Gard/1f9ee209316a09ba.jpg')
```

```
img_path: images/test/26.Pont_du_Gard/1f9ee209316a09ba.jpg
{'position': 0, 'index': 26, 'value': 19.572773, 'label': '26.Pont_du_Gard'}
{'position': 1, 'index': 48, 'value': 9.511307, 'label': '48.Whitby_Abbey'}
{'position': 2, 'index': 38, 'value': 9.128225, 'label': '38.Forth_Bridge'}
```



```
In [ ]: suggest_locations('images/test/36.Badlands_National_Park/0b82cd7b4b0a16d7.jpg')
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```
img_path: images/test/36.Badlands_National_Park/0b82cd7b4b0a16d7.jpg
{'position': 0, 'index': 18, 'value': 12.210916, 'label': '18.Delicate_Arch'}
{'position': 1, 'index': 36, 'value': 11.308897, 'label': '36.Badlands_National_Park'}
{'position': 2, 'index': 42, 'value': 8.565559, 'label': '42.Death_Valley_National_Par
k'}
```

